


PATH PLANNING FOR TREASURE HUNTING WITH ADVERSARIES AND ITEMS

By Anand Patel
ENAE788V
Spring 2019


BACKGROUND

Design & solve a treasure hunting game with adversaries, items, and exit door goal.

- Games provide a complex, interesting testbed for motion planning
 - Unique rules
 - Actions alter environment
- Motivation:
 - Pursuit/Evasion Motion Planning Problems
 - Applications for entertainment
 - Human vs Robot: Can we beat the machines?
 - Robot vs Robot: Who designed better agent?



https://www.retrogamer.net/retro_games80/the-making-of-ms-pac-man/




<https://www.usgamer.net/articles/long-time-coming-finishing-the-original-legend-of-zelda-in-2016>

CONCEPT OVERVIEW

Game Description:

- 2D static grid \rightarrow Undirected Graph: $G \in (V, E)$
 - Rooms \rightarrow Nodes: $v \in V = [v_1, v_2, \dots, v_n]$
 - Connections between rooms \rightarrow Edges: $(v_i, v_j) \in E \Leftrightarrow (v_j, v_i) \in E$
 - 4 directions of movement, Turn-based
- Explorer Agent R
 - R begins game at fixed start node v_{start}
- Adversaries A move randomly
 - 3 adversaries (A_1, A_2, A_3) begin game at exit door D
- WIN: R reaches fixed exit door D
- LOSE: R, A occupy same node
- Loot L and Items I placed at unique nodes
 - Collected by R by traveling onto node
 - Total Loot value: $L_T = 103$
 - Loot collected by explorer: L_C




$L\text{-type}$	#	Value
L_1	1	13
L_2	2	11
L_3	3	7
L_4	4	5
L_5	5	3
L_6	6	2

PROBLEM STATEMENT

Problem: *Path Planning for Treasure Hunting with Adversaries and Items*


For a given labyrinth represented by G , containing explorer R , loot L , adversaries A , and items I , the explorer shall determine the path from the initial location of R to door D that maximizes the percentage of loot value collected $\frac{L_C}{L_T}$ while avoiding collision with any adversaries A .

→ Algorithm Chosen: Q-Learning Explorer Agent



ITEMS

- Items are used upon collection by explorer
- Item 1: I_1 = Charm, 9 turns
 - Nearest adversary: harmless to explorer
- Item 2: I_2 = Root, 5 turns
 - All adversaries: Stationary, but harmful
- Item 3: I_3 = Lightning, permanent
 - Kills nearest adversary: harmless & stationary



I -type	#	Type
I_1	2	Charm
I_2	2	Root
I_3	1	Lightning

ADVERSARY BEHAVIOR

Each turn...

Random Walk:


If in hallway:

Continue in same direction of movement


If @ intersection:

Randomly pick from possible directions at current node

■ *Non-Deterministic*



Blue circles = intersection nodes



Q-LEARNING

$Q(s,A)$ = Q-Table (multi-dimensional array), initialized @ 0

$$Q^{new}(s_t, a_t) = Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a_t) - Q(s_t, a_t))$$

Q-learning: An off-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

until S is terminal

Epsilon-greedy annealing:

For current observation S , select random action A with probability *epsilon*.

Otherwise choose action based on greatest Q-value (greedy):

$$A = \max_A Q(S, A)$$

PARAMETERS & REWARDS USED

- Learning Rate $\alpha = 0.0003$
[Segarro2017mspacman]
- Discount Factor $\gamma = 0.945$
[Segarro2017mspacman]
 - α & γ need trial and error for fine tuning
- Training = 7000 episodes, Testing = 500 episodes
- *Epsilon* decays per episode
 - 0.9 → 0.1
 - Decay rate chosen based on 7000 training episodes

EVENT	REWARD	DESCRIPTION
Loot	3*LootValue	Explorer enters loot room
Item	6	Explorer enters item room
Door	400	Explorer enters exit door room
Enemy	-1000	Explorer is in same room as hostile enemy
Step	-1	Explorer performs a move

```
% number of episodes for training on or testing on.  
numEpisodes = 7000;
```

```
if(Training == 1)  
    % Learning Rate  
    alpha = 0.0003;  
    % Discount Factor  
    gamma = 0.945;  
    % Annealing Exploration  
    epsilon = 0.9;    % starting epsilon value  
    epsilon_decay_rate = 0.0005;  
    epsilon_min = 0.1;    % minimum epsilon value  
elseif(Testing == 1)  
    alpha = 0;  
    gamma = 0;  
    % No exploration during testing  
    epsilon = 0;    % starting epsilon value  
    epsilon_decay_rate = 0;  
    epsilon_min = 0.1;    % minimum epsilon value  
end
```


REPRESENTING THE GAME ENVIRONMENT

- Avoid over-fitted exploring agent by representing states S with info that translates between G (intersection nodes, relative distances, etc.)


$$Q(S,A) = Q(A, a(A), b, c, d, e(A), f(A), g(A), h(A), i(A))$$

- Actions + 9 Higher-order state/action inputs = 10-D Q-Table
 - Observe S in Q-Learning → Calculate (S,A) for all actions available to the Explorer @ current node

STATE/ACTION INPUTS

$$Q(S,A) = Q(A, \overset{\downarrow}{a(A)}, \overset{\downarrow}{b}, \overset{\downarrow}{c}, \overset{\downarrow}{d}, e(A), f(A), g(A), h(A), i(A))$$

- **A**: [down, left, right, up] = [1,2,3,4]
- **a(A)**: DoorInput(A) = distance to door [0,24]
 - Manhattan Distance used for all distances
 - 24 is max distance between nodes in this G
- **b**: LightningInput = if lightning has been used [0 or 1]
- **c**: CharmInput = current duration of charm [0,9]
- **d**: RootInput = current duration of root [0,5]



STATE/ACTION INPUTS

$$Q(S,A) = Q(A, a(A), b, c, d, e(A), f(A), g(A), h(A), i(A))$$

$e(A)$: EnemyInput(A) = $(D_{max} - D_{i,explorer} + D_{i,adversary}) [0,24]$

D_{max} : 24


$D_{i,explorer}$: distance between explorer and closest intersection along action A

$D_{i,adversary}$: distance between closest intersection along A and the nearest adversary

→ Input is capped at 24; equidistant is worst case threshold

Gives adversary threat for a given action A

[bom2013]



Why nearest enemy only considered.

For the “left” action, even though 2 ghosts approach, only the closest one to Ms. Pacman determines danger since they share the worst input value.

[bom2013]



How input changes.

For the “right” action, the ghost being closer to the intersection than Ms. Pac-Man makes this the highest threat level (input = 24). If the ghost was further away from the intersection than Ms. Pacman, the input would be lower.


STATE/ACTION INPUTS

$$Q(S,A) = Q(A, a(A), b, c, d, e(A), f(A), g(A), h(A), i(A))$$

- $f(A)$: LootProxInput(A) = distance to nearest loot [0,24]
 - Tells which action gets explorer closer to loot
- $g(A)$: ItemProxInput(A) = distance to nearest item [0,24]
 - Tells which action gets explorer closer to item

STATE/ACTION INPUTS

[bom2013]




Shows entrapped Ms.Pacman approaching SafeRoute 1 intersection away as best chance for survival

$$Q(S,A) = Q(A, a(A), b, c, d, e(A), f(A), g(A), h(A), i(A))$$

- $h(A)$: $\text{EntrapmentInput}(A) = (\text{SafeRoutes} - \text{SafeRoutes}(A)) / \text{SafeRoutes}$ [0,10]
 - Safe route = Shortest path from explorer to 3 intersections away that can be reached before any adversary
 - SafeRoutes = number of safe routes from current node
 - $\text{SafeRoutes}(A)$ = number of safe routes beginning with action A
 - If $\text{length}(\text{SafeRoutes}) = 0$, use 2 intersections away, 1 intersection away etc.
 - If 1 intersection away = 0, set input to worst case $\rightarrow 10$ (totally entrapped)
 - % of safe routes NOT along an action A
-
- $i(A)$: $\text{SameDirectionInput}(A) =$ if explorer is traveling in same direction [0 or 1]

TRAINING (IN PROGRESS)

- Training on 7000 episodes [bom2013], [segarro2017]
 - Started Monday night, still training
 - Fixed loot/item locations
 - No graphics → visualization done post-training/testing



Training trial for 10 episodes:

Explorer (red) uses 1xLightning, 1xCharm.

Killed adversary (yellow) @ node 68.

LOSS: caught by adversary (green) @ node 43.

TESTING (NEXT STEPS)

When training completes...

- Testing on 500 episodes [segarro2017]
 - On same G (fixed loot/items)
 - On G with random loot/item locations per episode ← assess policy transfer
- Metrics to analyze:
 - Win %
 - % Loot Value collected per win
 - Items used per win ← impact of items on success

REFERENCES

- [1] A. D. Tijsma, M. M. Drugan, and M. A. Wiering, “Comparing exploration strategies for q-learning in random stochastic mazes,” in *2016 IEEE Symposium Series on Computational Intelligence (SSCI)*. IEEE, 2016, pp. 1–8.
- [2] R. A. Bianchi, C. H. Ribeiro, and A. H. Costa, “Heuristically accelerated q-learning: a new approach to speed up reinforcement learning,” in *Brazilian Symposium on Artificial Intelligence*. Springer, 2004, pp. 245–254.
- [3] L. L. DeLooze and W. R. Viner, “Fuzzy q-learning in a nondeterministic environment: developing an intelligent ms. pac-man agent,” in *2009 IEEE Symposium on Computational Intelligence and Games*. IEEE, 2009, pp. 162–169.
- [4] M. Emilio, M. Moises, R. Gustavo, and S. Yago, “Pac-mant: Optimization based on ant colonies applied to developing an agent for ms. pac-man,” in *Proceedings of the 2010 IEEE Conference on Computational Intelligence and Games*. IEEE, 2010, pp. 458–464.
- [5] L. Bom, R. Henken, and M. Wiering, “Reinforcement learning to train ms. pac-man using higher-order action-relative inputs,” in *2013 IEEE Symposium on Adaptive Dynamic Programming and Reinforcement Learning (ADPRL)*. IEEE, 2013, pp. 156–163.
- [6] A. Segarra, “q-mspacman,” <https://github.com/albertsgrc/q-mspacman>, 2017.



QUESTIONS?